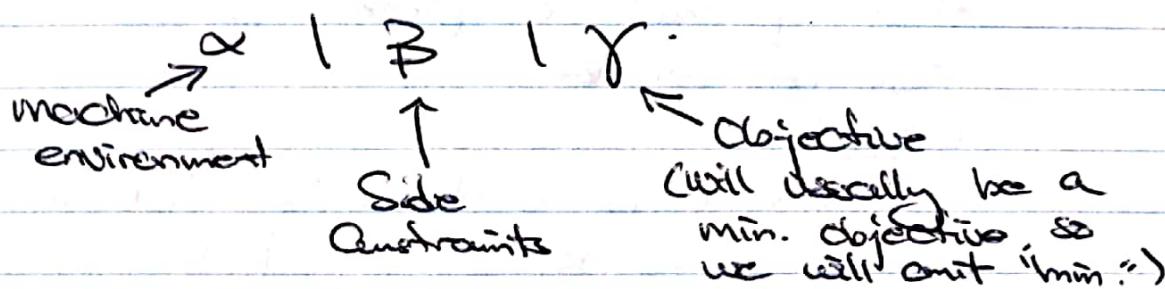


(4) Approximation algorithms

Notation:

We will refer to scheduling problems using triplet notation:



Ex:

α (M/C environment)

- 1': Denotes single M/C
- P: Multiple identical M/C's
- P_m or P_{m'}: m identical M/C's
(... and more to come)

γ (Objective Function)

- $\sum C_j$, $\sum w_j C_j$, Lmax, Fmax (Things we have seen)
- $\sum p_j(C_j)$, Conv_j := $\max_{i \in I} C_{ij}$ (Other possible objective functions)
Note: This is trivial with a single machine

β (Side constraints)

- This can be empty (which is exactly what we've seen so far in the B&B setup)
i.e. $\sum C_j$, $\sum w_j C_j$, Lmax, Fmax.
- There are other constraints we will see:
 - r_j : Release Dates
Specifies that job j becomes available at time r_j
i.e. Cannot schedule j at time $< r_j$)

↳ continued

Ex: (continued)

- Pre : Precedence Constraint

A precedence constraint $j \rightarrow k$ specifies that job j must complete before job k can be started.

Note:

Precedence is a transitive relation

i.e. $j \rightarrow k, k \rightarrow l \Rightarrow j \rightarrow l$.

This also implies that we cannot have a cycle of precedence constraints:

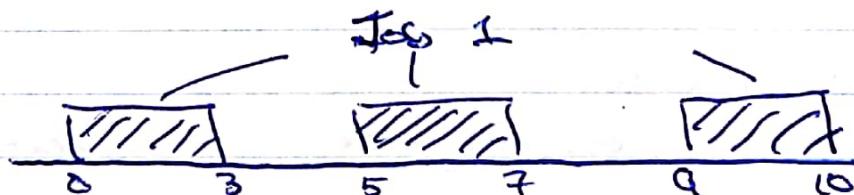
i.e. $j_1 \rightarrow j_2 \rightarrow \dots \rightarrow j_k \rightarrow j_1$ is
(not possible)

Precedence constraints are represented
by DAGs (i.e. Nodes := jobs, Arcs := Constraints)

- Pmt/Pmp : Preemption

Can interrupt a job and resuming
working on it later.

i.e. like our house!



as a valid schedule.

I|prec, fmax:

By our notation, this is the problem:

I | C_i, min. fmax, and subject
to precedence constraints $j \rightarrow k$.

We define:

Given a set J of jobs, let

$$\begin{aligned}L(J) &:= \{j \in J : j \text{ has no successors in } J\} \\&:= \{k \in J \text{ s.t. } j \rightarrow k\}\end{aligned}$$

Observe: Only jobs in L(J) can appear at the end of a schedule for J.

Def'n 4.1 (Modified LCT for I|prec, fmax)

- $J \leftarrow \{1, \dots, n\}$
- while $J \neq \emptyset$:
 - find $l \in L(J)$ s.t. $f_l(\sum_{j \in J} p_j) = \min_{k \in L(J)} f_k(\sum_{j \in J} p_j)$
 - Schedule l last
 - $J \leftarrow J \setminus \{l\}$ (and update L(J)).

Theorem 4.1:

The above algorithm gives an optimal schedule for I|prec, fmax.

[Proof]

Similar to the proof of Thm 2.4. □

$\hookrightarrow I(r_j, p_{min}, S_j)$

I(r_j , p_{j1}) \dots p_{jk})

We have:

I MC, min IC_j, and with the constraints of release dates and allowing preemption.

(Note: without r_j , preemption is unnecessary)

Def'n 4.2 (Shortest Remaining Processing Time Rule - SRPT)

At each point of time, schedule the job available with smallest remaining processing time, preempting a job if a job with smaller processing time is released

Theorem 4.2:

SRPT generates an optimal schedule for I(r_j , p_{j1}) \dots p_{jk} .

[Proof]

Let S^* be an optimal schedule that is NOT an SRPT schedule.

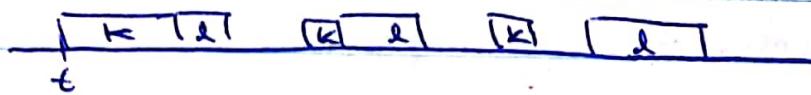
Then, at some point of time to S^* schedules some job k with remaining proc time x_k , but there was another available job l with remaining proc time $x_l < x_k$ and $x_l > 0$

Continued

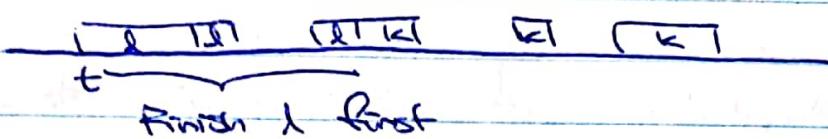
[Proof] (Continued)

S^* :

(Placement of k, l arbitrary)



S :



We want to be careful with our interchange argument (Else we might affect the completion time of other jobs)

Consider portions of time after t where k and l are scheduled in S^* , and interchange k, l by scheduling l in the first x_l of the time, followed by k to get schedule S .

This ensures that $C_j^S = C_j^{S^*} \quad \forall j \neq k, l$, and also $C_k^S = \max(C_k^{S^*}, C_l^{S^*})$.

Claim 1: $C_l^S < C_l^{S^*}$

Since k is processed for time $x_k > x_l$.

Claim 2: $C_l^S < C_k^{S^*}$

l is not scheduled at time t in S^* , so $C_l^{S^*}$ will be pushed back.

So, $C_l^S < \min(C_k^{S^*}, C_l^{S^*})$

$$\Rightarrow C_l^S + C_k^S < C_k^{S^*} + C_l^{S^*}$$

Completing the contradiction

Q.E.D

Hilary

$\max_j \sum C_j$:

Now, we look at the previous problem, but without preemption

Notice that SRPT generalizes SPT, so:

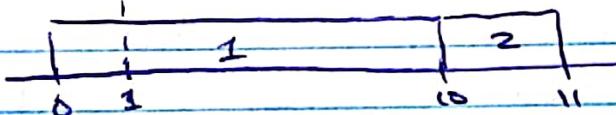
Q: Can we come up with generalizations or adaptations of SRPT/SPT to solve $\max_j \sum C_j$?

Attempt 1:

At each time t , schedule the available jobs with shortest processing time, but do not preempt.

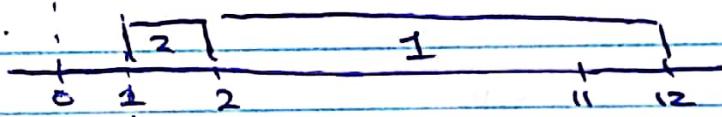
However, consider:

S:



$$\sum C_j^S = 10 + 11 \\ = 21$$

S^* :



$$\sum C_j^{S^*} = 2 + 12 \\ = 14$$

↑ ↑ $J_2 \leq J_1 : r_j = 1, p_j = 1$

~~$r_j=1$~~
 $r_j=0, p_j=10$

Notice that S^* is more optimal than S , so our algorithm is no longer optimal.

The problem is: A small job got released soon after a big job was started

→ Attempt 2

Attempt 2!

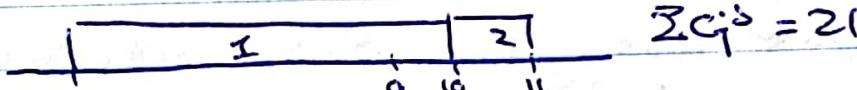
We might try to fix our previous example.
Consider:

Sort jobs in increasing P_j order, and process them in this order, waiting for job j , if j is not released by the completion time of the previous job.

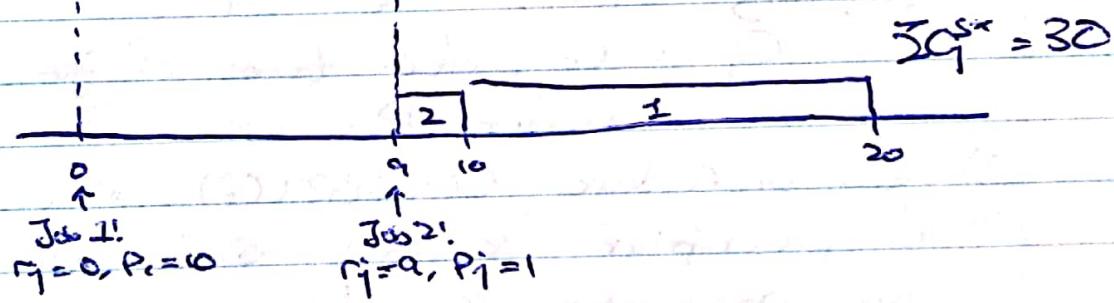
(Notice that this suggests that it may be beneficial to leave the machine idle even if there are available jobs - called "conceded idle time")

However, consider:

$S:$



$S^*:$



Again, our algorithm doesn't work since S is more optimal than S^* .

(An optimal solution to $\prod r_j \sum C_j$ is not known!)

What can we do?

- Abandon Efficiency
- Abandon Optimality.

↳ Continued

Hilroy

Def'n 4.3 (α -approximation algorithm)

An α -approximation algorithm is an efficient algorithm that returns a schedule S whose objective value $\leq \alpha \cdot (\text{optimum})$ (where optimum is the optimal objective value for our given problem, and we will commonly write it as OPT)

Further, α is known as the "approximation ratio" or "approximation guarantee".

We will design a 2-approximation algorithm for $\sum r_j | \sum C_j$:

- 1) Compute the preemptive schedule P by applying SPT, and let
 C_j^P = Completion time of job j under P .
- 2) Run procedure CONVERT(P), which takes a preemptive schedule and returns a non-preemptive schedule

CONVERT(P):

- 1) Sort jobs in increasing order of C_j^P :
 $C_1^P \leq C_2^P \leq \dots \leq C_n^P$
- 2) Schedule jobs "non-preemptively" in the order defined above, and output the resulting schedule
(By "non-preemptively", we mean:
Start Job j at time
 $t := \max(\text{completion time of } j-1, r_j)$).

Theorem 4.3:

SRPT + CONVERT is a 2-approx. algorithm for $\sum r_j \sum C_j$.

To prove this, we need to following:

Claim 4.4:

Let $OPT :=$ optimal value for $\sum r_j \sum C_j$.

Then:

$$OPT \geq \sum C_j^P$$

[Proof]

We've proved that SRPT gives an optimal schedule for $\sum r_j, \text{primal } \sum C_j$ (Theorem 4.2), so:

$$\sum C_j^P = OPT_{\sum r_j, \text{primal } \sum C_j}$$

And:

$$OPT_{\sum r_j, \text{primal } \sum C_j} \leq OPT$$

Since any schedule for $\sum r_j \sum C_j$ is feasible for $\sum r_j, \text{primal } \sum C_j$. The preemptive problem is a relaxation of $\sum r_j \sum C_j$. \square

Lemma 4.5:

For every job j (in the ordering after step (1) of CONVERT):

$$C_P \leq 2C_j^P$$

→ Proof

Hilary

[Proof]

We begin with the following claim:

For every job j (under the order after Step 5):

$$C_j^P \geq \max_{k=1, \dots, j} r_k, \text{ and } (1)$$

$$C_j^P \geq \sum_{k=1}^j p_k \quad (2)$$

Since by definition:

$$C_j^P \geq C_k^P \quad \forall k=1, \dots, j$$

$$\geq r_k \quad \forall k=1, \dots, j$$

$$\geq \max_{k=1, \dots, j} r_k$$

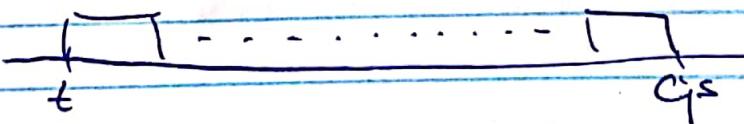
And, by time C_j^P , all jobs $1, \dots, j$ have completed in the preemptive schedule P

So,

$$C_j^P \geq \sum_{k=1}^j p_k$$

Now, in schedule S , let $t = \max_{k=1, \dots, j} r_k$, and
[] consider the interval $[t, C_j^S]$

S' :



Goal: We want to bound this interval by C_j^P

→ Continued

[Proof] (Continued)

Notice that:

- There is no idle time in $[t, C_j^S]$. Since the only way this can happen is if we are waiting for a job to be released. However, by time t , all jobs $1, \dots, j$ have been released and further, we would complete j before waiting for any other job i with $i > j$.
- This also means that in $[t, C_j^S]$, we are only processing jobs $1, \dots, j$.

Hence,

$$C_j^S = t + (\text{length of } [t, C_j^S])$$

$$\leq \max_{k=1, \dots, j} r_k + \sum_{k=1}^j p_k$$

(By the 2nd point)

$$\leq C_1^P + C_j^P = 2C_1^P$$

Note:

If we can solve a preemptive version of a non-preemptive problem, COVERT and this lemma gives a strong guarantee for the completion times.

Now, we can prove Theorem 4.3!

[Proof] (of Theorem 4.3):

By Lemma 4.5, we get that:

$$\sum_i C_i^S \leq 2 \sum_i C_i^P$$

And by claim 4.4:

$$\sum_i C_i^S \leq 2 \sum_i C_i^P \leq 2 \cdot OPT.$$

□

Hilroy

Runtime Analysis!

SRPT requires:

- $O(n^2)$:

- The schedule changes at most $2n$ time points corresponding to completion times + release dates of jobs

- At each of these time points, we go through the list of remaining jobs in $O(n)$ time

$$\text{So, } O(n \cdot 2n) = O(n^2)$$

- $O(n \log n)$

- By using a min-heap or a sorted list, we can reduce the $O(n)$ lookup to $O(\log n)$

$$\text{So, } O(n \log n)$$

Then, our 2-approx. alg takes:

SRPT + CONVRRT

\downarrow
 $O(n \log n)$

\downarrow
 $O(n \log n)$ (sorting)

= $O(n \log n)$ time.

$\text{ILr}_j(\sum w_j c_j)$ and $\text{ILr}_j(\text{prtnl } \sum w_j c_j)$

We know that $\text{ILr}_j(\text{prtnl } \sum w_j c_j)$ is a relaxation of $\text{ILr}_j(\sum w_j c_j)$ and we have a general strategy to tackle preemptive problems and convert them to their non-preemptive counterparts. So, we want an α -approx. algorithm A for $\text{ILr}_j(\text{prtnl } \sum w_j c_j)$ so that $A + \text{CONVERT}$ gives an 2α -approx algorithm for $\text{ILr}_j(\sum w_j c_j)$.

Preemptive WSPT:

At each point of time, schedule available jobs with largest w_j/p_j ratio, preempting as necessary.

Equivalently:

- 1) Order jobs in decreasing order of density
- 2) Schedule jobs, preemptively, in this order
(i.e. At each point of time, schedule job that came earliest in this order)

Note: Preemptive WSPT does not coincide with SPT, even when all jobs have weight 1, since p_j is the original processing time for j .

Theorem 4.6:

Preemptive WSPT is a 2-approx. algorithm for $\text{ILr}_j(\text{prtnl } \sum w_j c_j)$.

→ proof.

Hilroy

[Proof]

Let $\text{OPT} := \text{OPT}_{\{1|r_j, p_m| \sum w_j C_j\}}$

let jobs be ordered $1, \dots, n$ so that:

$$\frac{w_1}{p_1} \geq \frac{w_2}{p_2} \geq \dots \geq \frac{w_n}{p_n}$$

and S be the schedule returned by
preemptive WSPT

(1)
Observe that $\text{OPT} \geq \sum_{j=1}^n w_j C_j$. Since the

completion time $C_j \geq r_j$ in any schedule.

And, $\text{OPT} \geq \text{OPT}_{\{1|p_m| \sum w_j C_j\}}$, which is the
relaxation of the problem, with all
release dates at 0. (Any schedule for

$\{1|r_j, p_m| \sum w_j C_j\}$ is a schedule $\{1|p_m| \sum w_j C_j\}$)

Further $\{1|p_m| \sum w_j C_j\} = \{1| \sum w_j C_j\}$, since
preemption is unnecessary when $r_j = 0 \forall j$,
so we can use WSPT, which produces
an optimal schedule with objective value
 $\text{OPT}_{\{1| \sum w_j C_j\}}$. The schedule will also be
exactly in the order $1, \dots, n$, so:

$$\text{OPT} \geq \text{OPT}_{\{1|r_j, p_m| \sum w_j C_j\}}$$

$$= \text{OPT}_{\{1| \sum w_j C_j\}}$$

$$= \sum_{j=1}^n w_j \left(\sum_{k=1}^j p_k \right). \quad (2)$$

Continued

[Proof] (Continued)

We are nearly done. Consider schedule S and job j .

$S:$



In the interval $[r_j, C_j^S]$, there is never any idle time, since job $\neq j$ is available. And, only jobs k ($1 \leq k \leq j$) can be scheduled. Since any job k , with $k \geq j+1$ would contradict our ordering.

So:

$$C_j^S \leq r_j + \sum_{k=1}^j p_k \quad \forall j = 1, \dots, n$$

Then:

$$\begin{aligned} \sum_{j=1}^n w_j C_j^S &\leq \sum_{j=1}^n w_j (r_j + \sum_{k=1}^j p_k) \quad (\text{By above}) \\ &= \sum_{j=1}^n w_j r_j + \sum_{j=1}^n w_j \left(\sum_{k=1}^j p_k \right) \\ &\leq OPT_1 + OPT_2 = 2 \cdot OPT \end{aligned}$$

II.

And so, Preemptive WSPT \rightarrow CONVERT gives an algorithm for $\sum w_j C_j$.

Corollary 4.7:

Preemptive WSPT \rightarrow CONVERT is a 4-approx. alg for $\sum w_j C_j$